

Deep Learning

Big Data & Machine Learning Bootcamp - Keep Coding







Outline

1. Deep Neural Network
2. Why deep representations?
3. Regularization
4. Normalizing Inputs
5. Vanishing/Exploding gradients



Deep Neural Network

“Shallow network”



Andrew Ng

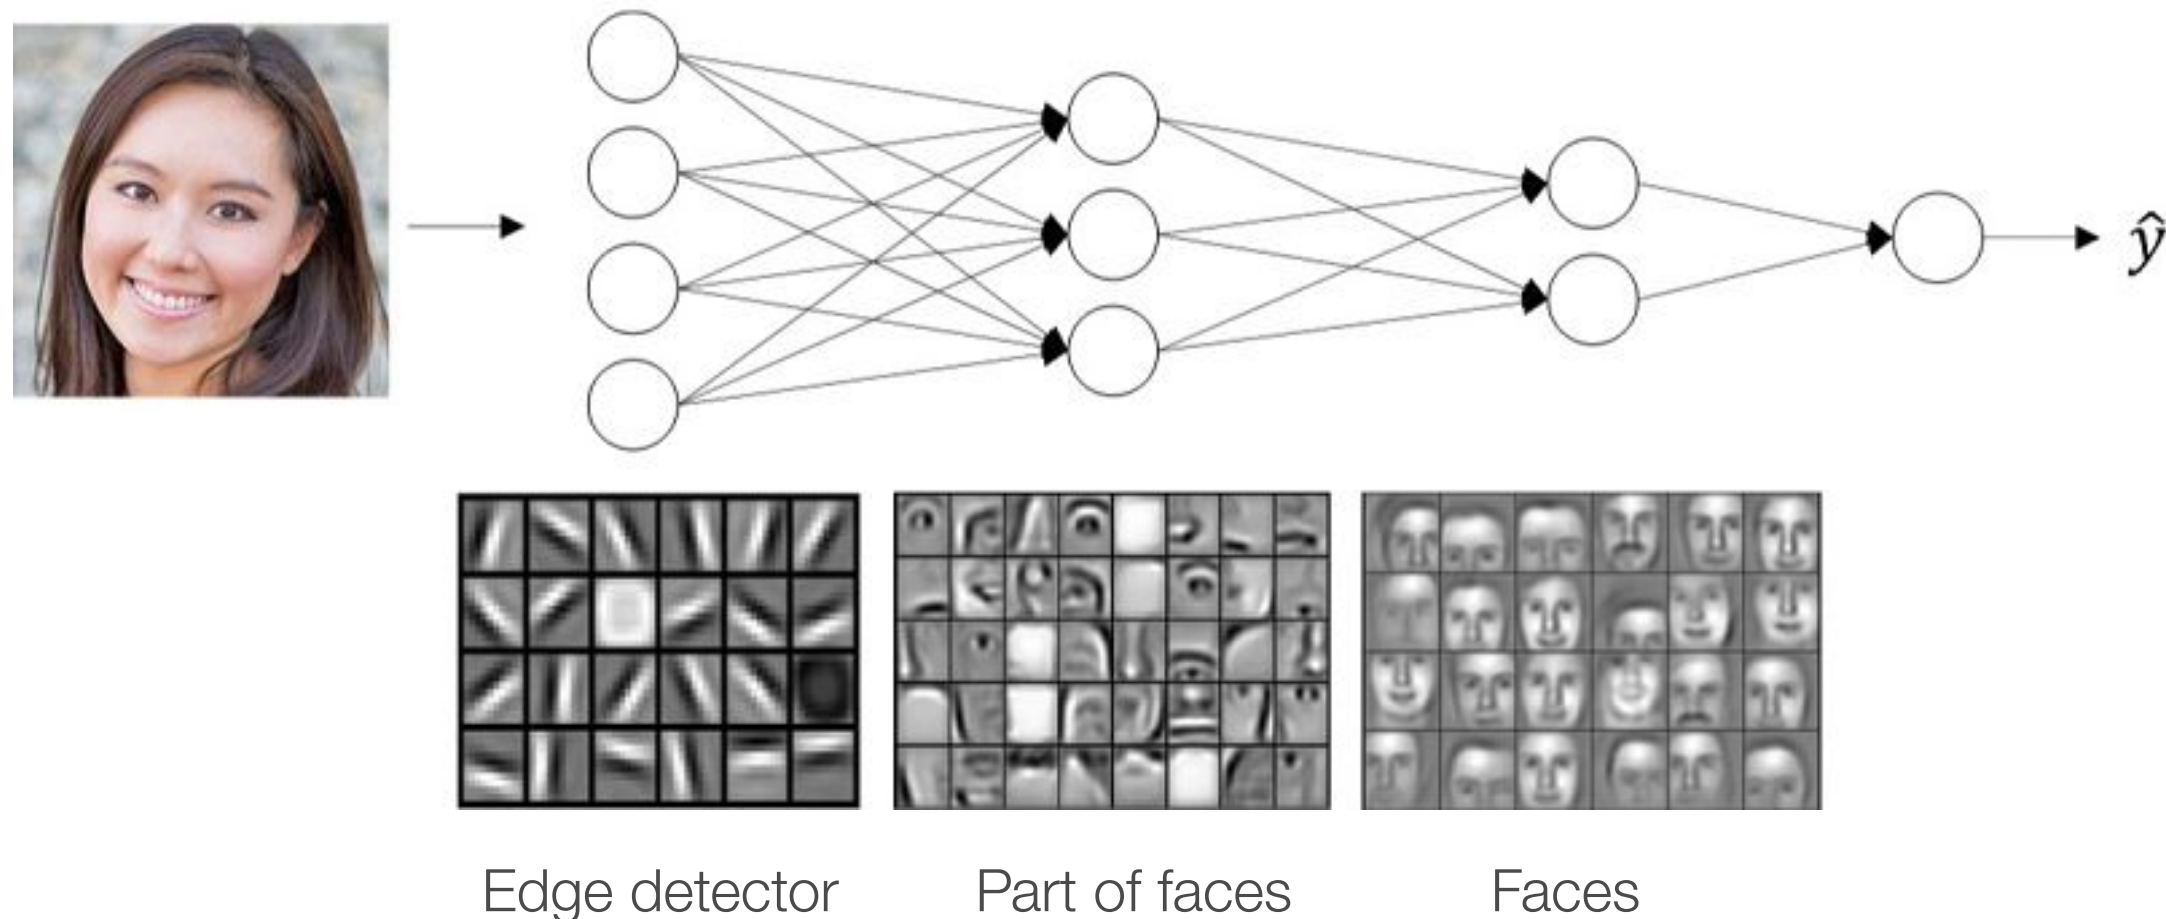
We don't count the input and the output layers. We only count the hidden layers!

Sources:
- Coursera




Why deep representations?

The first layers detect low level features while deep layers detect more complex features





Sources:
- Coursera

Bias/variance trade-off



Underfitting



Overfitting



Sources:
- Coursera

© All rights reserved. www.keepcoding.io

Train - Development and Test sets

Train set error	1%	15%	15%	0.5%
Dev set error	11%	16%	30%	1%
	High variance (Overfitting)	High bias (Underfitting)	High bias & high variance	Low bias and low variance

This is assuming that the human error or the optimal error is 0%

*We'll discuss this later. But this doesn't occur in real life. **The optimal error is around 15%***



Sources:
- Coursera

© All rights reserved. www.keepcoding.io

Regularization

L2 regularization: Add a term to the loss function that is function of the parameters in the neural network


L2 regularization is also called **weight decay**.




Sources:
- Coursera

Regularization

Dropout is also another regularization technique! How does it work?



(a) Standard Neural Network



(b) Network after Dropout

*Essentially you define a **probability of keeping a neuron.***

A common way of implementing dropout is using the “inverted dropout”



Regularization

There are other regularization techniques such as:

- **Data augmentation** (mirroring, horizontal and vertical rotation, zooming, etc)
- **Early stopping:** You stop the training process when the dev set error gets bigger when compared to the training set error.

Early stopping is **not** recommended as it breaks orthogonalization. This means, orthogonalization doesn't allow you to work on optimizing the cost function and avoiding overfitting independently.





Sources:
- Coursera



Normalizing Inputs

The main reason for normalizing the inputs is that the cost function is more rounded and easy to optimize!

Unnormalized:



Normalized:



The main idea is to keep the similar ranges among the input variables



Sources:
- Coursera

Vanishing/Exploding gradients

When training very deep neural networks, **slopes or the derivatives can get very big (exploding) or very very small (vanishing)**

To remember:

- Weight values that are smaller than 1 make the updates to decrease when doing multiplication (***vanishing gradients***)
- The contrary happens when having derivative values bigger than 1, updates increase (***exploding gradients***)

A proper weight initialization method may help to reduce this problem!

